

## HW 9

### Problem 1

```
x <- c(-.86, -.3, -.05, .73)
n <- rep(5, 4)
y <- c(0, 1, 3, 5)
```

Checking the gradient.

```
grad <- function(a, b) {
  c(sum(y/(1+exp(a+b*x)) - (n-y)/(1+exp(-a-b*x))),
    sum(x*y/(1+exp(a+b*x)) - x*(n-y)/(1+exp(-a-b*x))))
}
```

The analytic gradient evaluated at (1,10) is

```
grad(1, 10)
```

```
## [1] 0.2904300 -0.1125231
```

The finite-difference approximation of gradient at (1,10) is

```
logp <- function(a, b) {
  sum(-y*log(1+exp(-a-b*x)) - (n-y)*log(1+exp(a+b*x)))
}
```

```
(logp(1+1e-5, 10) - logp(1, 10))/1e-5
```

```
## [1] 0.2904215
```

```
(logp(1, 10+1e-5) - logp(1, 10))/1e-5
```

```
## [1] -0.1125233
```

They do match up to several decimal places.

### Implementation

```
set.seed(0)

hmc <- function(iter, init) {
  Theta <- matrix(0, nrow=iter+1, ncol=2)
  Theta[1, ] <- init

  M <- matrix(c(5, 0, 0, 5), 2, 2)
  eps <- .12
  L <- 10
  Accept <- rep(NA, iter)
```

```

for (i in 1:iter) {
  # Leapfrog
  Theta0 <- Theta[i, ]
  phi <- rnorm(2, mean=0, sd=diag(M))
  phi0 <- phi
  phi <- phi + .5*eps*grad(Theta0[1], Theta0[2])
  Theta0 <- Theta0 + eps*t(phi)%*%solve(M)
  for (j in 1:(L-1)) {
    phi <- phi + eps*grad(Theta0[1], Theta0[2])
    Theta0 <- Theta0 + eps*t(phi)%*%solve(M)
  }
  phi <- phi + .5*eps*grad(Theta0[1], Theta0[2])

  #Accept-reject sampling
  log_r <- logp(Theta0[1], Theta0[2]) - logp(Theta[i, 1], Theta[i, 2]) +
    sum(dnorm(phi, sd=diag(M), log=T)) - sum(dnorm(phi0, sd=diag(M), log=T))
  Accept[i] <- (log(runif(1)) < log_r)
  Theta[i+1, ] <- ifelse(rep(Accept[i], 2),
                        Theta0,
                        Theta[i, ])
}
return(list(chain=Theta, accept=Accept))
}

mean(hmc(1e4, c(0, 10))$accept)

## [1] 0.6668

```

## Effective sample size.

```

burn_in <- function(chain) {
  tail(chain, -.3*length(chain))
}

my_diff <- function(t, x) {
  mean(diff(x, lag = t)^2, na.rm = TRUE)
}

chain1 <- hmc(3000, c(0, 10))$chain
chain2 <- hmc(3000, c(10, 20))$chain
chain3 <- hmc(3000, c(-10, 30))$chain
chain4 <- hmc(3000, c(-20, 50))$chain

chains <- cbind(chain1[, 1], chain2[, 1], chain3[, 1], chain4[, 1])
chains <- apply(chains, 2, burn_in)

n_eff <- function(chains) {
  require(purrr)
  m <- ncol(chains)
  n <- nrow(chains)
  # Within-chain var
  W <- 1/m*sum(apply(chains, 2, var))
}

```

```

means <- colMeans(chains)
B <- n * var(means)

V_plus <- (n-1)/n * W + 1/n * B

V <- sapply(seq_len(nrow(chains) - 1), my_diff, x=chains[, 1]) + sapply(seq_len(nrow(chains) - 1), my_diff, x=chains[, 2])

V <- V/m

rho <- head_while(1 - V / (2*V_plus), ~. > 0)

n_hat_eff <- m*n / (1 + 2*sum(rho))
return(n_hat_eff)
}

n_eff(chains)

```

## Loading required package: purrr

## [1] 106.1946

Using my tuned hyper-parameters, it takes 3000 iterations to achieve 100 effective sample size.

## Check with direct approach

```
colMeans(chain1)
```

## [1] 1.448485 12.267518

Back in homework 3 the mean estimation using direct approach is (0.9960628, 10.5482460). It is consistent with the direct approach. There is error because the sample size is not large enough and we don't expect the estimation to be accurately matched.

## Problem 2

### Sampling and diagnostic

```

y <- c(53, 57, 66, 67, 72)
joint_density <- function(N, theta, y) {
  # y can be a vector
  prod(choose(N, y)*(theta^y)*(1-theta)^(N-y))
}

banana <- function(iter, init) {
  res <- matrix(nrow=iter+1, ncol=2)
  res[1, ] <- init #c(2*max(y), 1/2) init
  for (i in 1:iter) {
    temp_N <- rpois(1, lambda=res[i, 1]) # proposal
    log_ratio <- log(joint_density(temp_N, res[i, 2], y)) -
      log(joint_density(res[i, 1], res[i, 2], y)) +
      log(dpois(temp_N, lambda=res[i, 1])) -
      log(dpois(res[i, 1], lambda=temp_N))
    temp_log_U <- log(runif(1))
    if (temp_log_U <= log_ratio) {

```

```

    res[i+1, 1] <- temp_N
  } else {
    res[i+1, 1] <- res[i, 1]
  }

  temp_theta <- runif(1, min=max(0, res[i, 2]-.2), max=min(1, res[i,2]+.2)) # proposal
  log_ratio <- log(joint_density(res[i+1, 1], temp_theta, y)) -
    log(joint_density(res[i+1, 1], res[i, 2], y))
  # proposal ratio is always one so dropped
  temp_log_U <- log(runif(1))
  if (temp_log_U <= log_ratio) {
    res[i+1, 2] <- temp_theta
  } else {
    res[i+1, 2] <- res[i, 2]
  }
}
}
return(res)
}

```

```

chain1 <- banana(2e5, c(200, .5))
chain2 <- banana(2e5, c(100, .9))
chain3 <- banana(2e5, c(200, .1))
chain4 <- banana(2e5, c(100, .5))


```

It takes too long to compute the  $\hat{\rho}_t$  so I will do diagnostics only monitoring the mixing.


```

plot(chain1[, 1], type='l', main = 'N')
lines(chain2[, 1], col='red')
lines(chain3[, 1], col='green')
lines(chain4[, 1], col='blue')

```



```
plot(chain1[, 2], type='l', main = 'theta')
lines(chain2[, 2], col='red')
lines(chain3[, 2], col='green')
lines(chain4[, 2], col='blue')
```



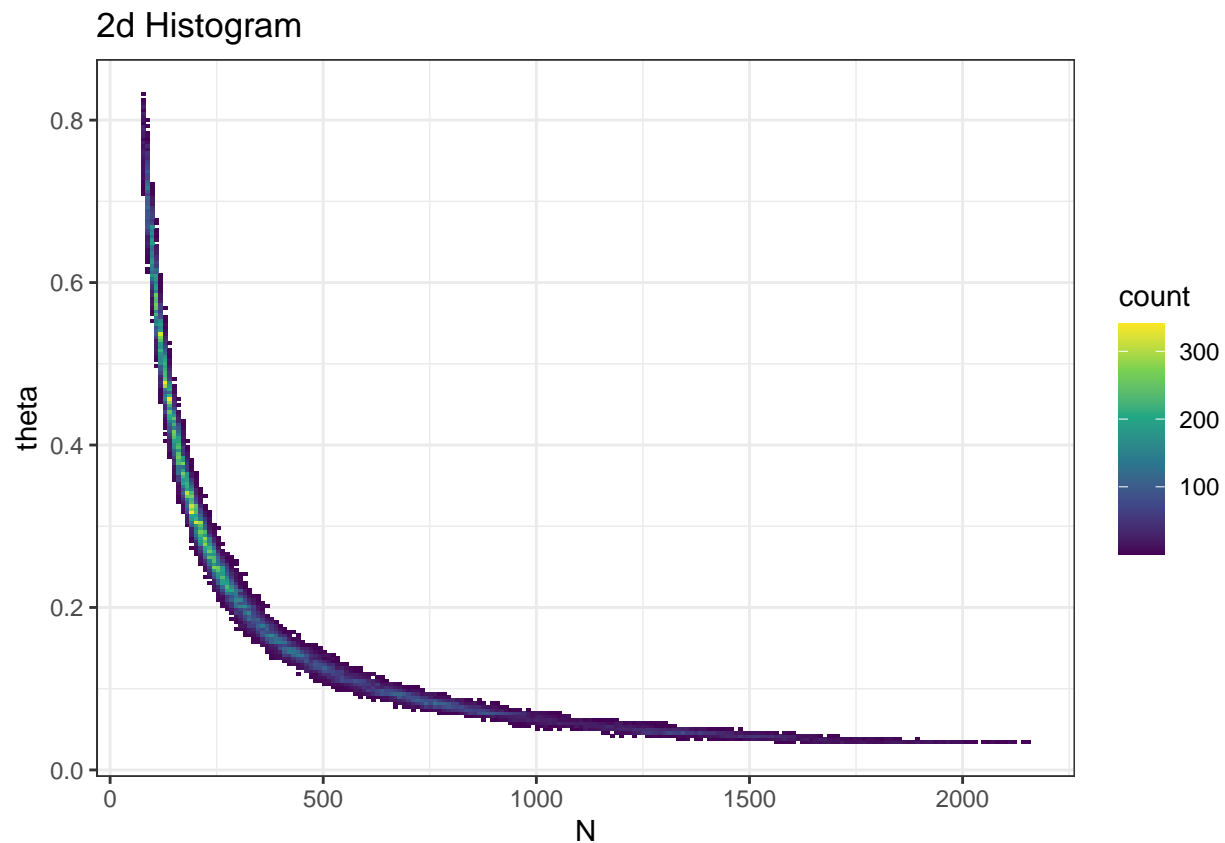
It looks like for  $\theta$  the chains are indistinguishable but for  $N$  there are some disturbance. I think it is good enough.

## Posterior density visualization

### 2d histogram

Let's first take a look at a "nice" plot. It looks better comparing to the scatterplot and contour plot.

```
library(ggplot2)
res <- burn_in(chain1) # remove burn-in
plt_data <- as.data.frame(res)
colnames(plt_data) <- c('N', 'theta')
ggplot(plt_data, aes(x=N, y=theta) ) +
  geom_bin2d(bins = 200) +
  scale_fill_continuous(type = "viridis") +
  theme_bw() + ggtitle('2d Histogram')
```




### Scatterplot and contour plot

It is not reasonable to do a 200000 size scatterplot so I just randomly picked 200 points from the chain to draw the scatterplot. The contour plot is using the entire chain.


```
library(MASS)
res <- burn_in(chain1)
ind <- sample(nrow(res), 200)
plot(res[ind, 1], res[ind, 2], col='blue',
      xlab='N', ylab='theta', main='Chain1: Scatter plot of 200 points & Contour plot')
contour(kde2d(res[, 1], res[, 2], n=100), add=TRUE)
```

### Chain1: Scatter plot of 200 points & Contour plot




```
res <- burn_in(chain2)
ind <- sample(nrow(res), 200)
plot(res[ind, 1], res[ind, 2], col='blue',
      xlab='N', ylab='theta', main='Chain2: Scatter plot of 200 points & Contour plot')
contour(kde2d(res[, 1], res[, 2], n=100), add=TRUE)
```

### Chain2: Scatter plot of 200 points & Contour plot



```
res <- burn_in(chain3)
ind <- sample(nrow(res), 200)
plot(res[ind, 1], res[ind, 2], col='blue',
      xlab='N', ylab='theta', main='Chain3: Scatter plot of 200 points & Contour plot')
contour(kde2d(res[, 1], res[, 2], n=100), add=TRUE)
```


### Chain3: Scatter plot of 200 points & Contour plot



```
res <- burn_in(chain4)
ind <- sample(nrow(res), 200)
plot(res[ind, 1], res[ind, 2], col='blue',
      xlab='N', ylab='theta', main='Chain4: Scatter plot of 200 points & Contour plot')
contour(kde2d(res[, 1], res[, 2], n=100), add=TRUE)
```



### Chain4: Scatter plot of 200 points & Contour plot



Posterior prob. estimation for event  $N > 100$ .

```
mean(res[, 1] > 100)
```

```
## [1] 0.9357883
```