

HW 10

Problem 1

(1) & (2)

We will skip the single run in (1) and do 100 repeat runs directly.

```
library(rstan)
options(mc.cores = parallel::detectCores())
rstan_options(auto_write = TRUE)

lr <- stan_model(file = 'LR.stan') # compile the model

beta <- c(3, .1, .5)
coverage <- matrix(NA, nrow=100, ncol=3)

for (i in 1:100) {
  x1 <- seq(1, 100, 1)
  x2 <- rbinom(100, 1, .5)
  err <- rt(100, df=4) * 5

  y <- 3 + .1*x1 + .5*x2 + err

  lr_dat <- list(N=100, x1=x1, x2=x2, y=y)

  chain <- sampling(object = lr, data = lr_dat, seed = 0)
  stats <- summary(chain)$summary
  coverage[i, ] <- (stats[1:3, 5] < beta) & (stats[1:3, 7] > beta)
}

colMeans(coverage, na.rm = T)

## [1] 0.47 0.53 0.53
```

(3)

```
lr_t <- stan_model(file = 'LR_t.stan') # compile the model

beta <- c(3, .1, .5)
coverage_t <- matrix(NA, nrow=100, ncol=3)

for (i in 1:100) {
  x1 <- seq(1, 100, 1)
  x2 <- rbinom(100, 1, .5)
  err <- rt(100, df=4) * 5

  y <- 3 + .1*x1 + .5*x2 + err
```

```

lr_dat <- list(N=100, x1=x1, x2=x2, y=y)

chain <- sampling(object = lr_t, data = lr_dat, seed = 0)
stats <- summary(chain)$summary
coverage_t[i, ] <- (stats[1:3, 5] < beta) & (stats[1:3, 7] > beta)
if (i/10 == round(i/10)) print(i)
}

## [1] 10
## [1] 20
## [1] 30
## [1] 40
## [1] 50
## [1] 60
## [1] 70
## [1] 80
## [1] 90
## [1] 100

colMeans(coverage_t, na.rm = T)

## [1] 0.52 0.53 0.55

```

Problem 2

Denote the number of shock avoidances made by dog j before trial i as W_{ij} and whether the dog j received the shock in trail j as $S_{ij} \in \{0, 1\}$ where 0 means no shock. Let T_{ij} be a latent variable indicating the waiting time for dog j in trial i to jump when presented the CS.

The model is $S_{ij} = I(T_{ij} \leq 10)$ where $T_{ij} \sim \text{Exp}(\lambda_{ij})$ and $\log \lambda_{ij} = \alpha i + \beta W_{ij}$. The prior distribution for α and β are both uniform distribution on $(0, 100)$. Please note this is a weakly informative prior because we are enforcing that the training effect is positive and also the effect is not dramatically strong.

The posterior is

$$\begin{aligned}
 p(\alpha, \beta, T|S) &\propto p(\alpha, \beta)p(T|\alpha, \beta)p(S|T) \\
 &\propto (\alpha i + \beta W_{ij}) \exp\{-(\alpha i + \beta W_{ij})T_{ij}\} I\{T_{ij}S_{ij} \in \{0\} \cup (10, +\infty)\}.
 \end{aligned}$$

```

dogs_model <- stan_model(file = 'Dogs.stan') # compile the model

dogs <- read.table('dogs.txt', header=F, fill=T)[-c(1:2), -1]
dogs[dogs == 'S'] <- 1
dogs[dogs == '.'] <- 0
W <- t(apply(dogs, 1, cumsum))
IND <- t(apply(matrix(1, 30, 25), 1, cumsum))

dogs_dat <- list(J=30, I=25, S=dogs, W=W, IND=IND)
chain <- sampling(object = dogs_model, data = dogs_dat,
                  seed = 0, iter=5000, chains=1)

##
## SAMPLING FOR MODEL 'Dogs' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.000245 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 2.45 seconds.

```

```
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration: 1 / 5000 [ 0%] (Warmup)
## Chain 1: Iteration: 500 / 5000 [ 10%] (Warmup)
## Chain 1: Iteration: 1000 / 5000 [ 20%] (Warmup)
## Chain 1: Iteration: 1500 / 5000 [ 30%] (Warmup)
## Chain 1: Iteration: 2000 / 5000 [ 40%] (Warmup)
## Chain 1: Iteration: 2500 / 5000 [ 50%] (Warmup)
## Chain 1: Iteration: 2501 / 5000 [ 50%] (Sampling)
## Chain 1: Iteration: 3000 / 5000 [ 60%] (Sampling)
## Chain 1: Iteration: 3500 / 5000 [ 70%] (Sampling)
## Chain 1: Iteration: 4000 / 5000 [ 80%] (Sampling)
## Chain 1: Iteration: 4500 / 5000 [ 90%] (Sampling)
## Chain 1: Iteration: 5000 / 5000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 451.801 seconds (Warm-up)
## Chain 1: 455.362 seconds (Sampling)
## Chain 1: 907.163 seconds (Total)
## Chain 1:
```

```
stats <- summary(chain)$summary
stats[1:2, 9:10]
```

```
##          n_eff      Rhat
## alpha 12.569297 1.007875
## beta   9.935438 1.548232
```

For the summary we can tell how well it converged using `n_eff` and `Rhat`. Although the `n_eff` and `Rhat` are not satisfying, we will use 4000 iterations at this time because the program is running too slow for larger sample size.

Vectorization will help speed up but I could not find convenient way to vectorize this program and remove the for loop at this point.

Following is the posterior inference on (α, β) .

```
stats[1:2, c(1, 4:8)]
```

```
##          mean      2.5%      25%      50%      75%      97.5%
## alpha 0.020596677 0.020385227 0.020484875 0.020586324 0.020671442 0.02095566
## beta  0.009187564 0.007054933 0.008103831 0.009364539 0.009966576 0.01152375
```

Posterior inference on T_1 (mean and median only).

```
stats[3:27, c(1, 6)]
```

```
##          mean      50%
## T[1,1] 41.75953910 41.393792356
## T[1,2]  2.35453912  2.237662569
## T[1,3]  5.10452635  5.105817052
## T[1,4] 15.79432810 15.704713677
## T[1,5]  1.67372394  1.655419089
## T[1,6] 27.24510895 27.887315058
## T[1,7]  0.68327047  0.699448637
## T[1,8]  5.66276154  5.643688848
## T[1,9]  0.33140361  0.334146627
## T[1,10] 2.31671671 2.312742968
```

```
## T[1,11] 4.18057826 4.258493729
## T[1,12] 1.15175788 1.129763969
## T[1,13] 4.85814431 4.840310104
## T[1,14] 0.90584605 0.905724569
## T[1,15] 0.21650298 0.200198477
## T[1,16] 9.15181639 9.207317028
## T[1,17] 0.36949873 0.359432960
## T[1,18] 4.47524874 4.520395080
## T[1,19] 0.83457638 0.844777769
## T[1,20] 1.05374010 1.051074735
## T[1,21] 3.56121346 3.698212388
## T[1,22] 3.82027084 3.769465958
## T[1,23] 2.15468285 2.171327547
## T[1,24] 0.35080644 0.361198969
## T[1,25] 0.01013905 0.009890074
```

Problem 3

Since the total number of 3×3 binary matrices with grand total equals 4 is $\binom{9}{4=126}$, the proposal is a uniform distribution on $\{1, \dots, 126\}$ and our target is uniform on $\{1, \dots, 5\}$ which represent the five possible arrangements of matrices.

```
samp <- rep(0, 1e4)
accept <- c()
target <- seq(1, 5, 1)
i <- 1
while (i <= 1e4) {
  ind <- sample(1:choose(9, 4), 1)
  if (ind %in% target) {
    samp[i] <- ind
    accept <- c(accept, 1)
    i <- i + 1
  } else {
    accept <- c(accept, 0)
  }
}
mean(accept)
```

```
## [1] 0.03942487
```

The acceptance rate is expected to be low because we are trying to sample 5 out of 126 uniformly. And the experiment shows that the acceptance rate is indeed low.